

Large Language Models: Comprehensive Guide from Zero to Hero

A deep-dive into LLM architecture, training, fine-tuning, and deployment with in-depth explanations and working code examples.

Table of Contents

1. [Foundations of NLP and Machine Learning](#)
 2. [The Transformer Architecture: Core Concepts](#)
 3. [LLM Pre-training and Training Objectives](#)
 4. [Adaptation and Fine-Tuning](#)
 5. [Prompt Engineering and Prompting Techniques](#)
 6. [Retrieval-Augmented Generation and Agents](#)
 7. [Evaluation, Optimization, and Safety](#)
 8. [Deployment and Serving](#)
 9. [Training Your Model on Custom Data](#)
 10. [Replicating GPT-2 from Scratch](#)
 11. [Production LLM Systems](#)
-

1. Foundations of NLP and Machine Learning

Introduction to NLP: Text Preprocessing Pipeline

Natural Language Processing starts with converting raw text into structured data that neural networks can process. This pipeline is fundamental to all downstream tasks.

Tokenization: Breaking Text into Meaningful Units

Tokenization is the process of splitting text into tokens (words, subwords, or characters). This is more complex than simple whitespace splitting.

Example: Different tokenization approaches

```
text = "ChatGPT's capabilities are impressive!"
```

1. Word tokenization (naive)

```
word_tokens = text.lower().split()
```

Output: ["chatgpt's", "capabilities", "are", "impressive!"]

2. Subword tokenization (BPE - Byte Pair Encoding)

```
from tokenizers import Tokenizer
tokenizer = Tokenizer.from_file("gpt2.json")
tokens = tokenizer.encode(text)
```

Output: [17894, 10905, 1039, 2620, 28]

3. Character-level tokenization

```
char_tokens = list(text)
```

Output: ['C', 'h', 'a', 't', 'G', 'P', 'T', "", 's', ...]

Key Concepts:

- **Word-level tokenization:** Vocabulary size ~50,000-100,000. Simple but can't handle OOV (out-of-vocabulary) words
- **Subword tokenization (BPE/WordPiece):** Vocabulary size ~30,000-50,000. Modern standard for large models
- **Character-level:** Vocabulary size ~100-256. Can represent any word but sequence becomes very long

Stemming vs Lemmatization

Both reduce words to base forms, but with different approaches:

```
from nltk.stem import PorterStemmer, WordNetLemmatizer
```

```
stemmer = PorterStemmer()
lemmatizer = WordNetLemmatizer()
```

```
words = ["running", "runs", "ran", "easily", "fairly"]

print("Stemming:")
for word in words:
    print(f"{word} → {stemmer.stem(word)}")
```

Output:

running → run

runs → run

ran → ran (doesn't recognize past tense)

easily → easili

fairly → fairli

```
print("\nLemmatization:")
for word in words:
    print(f"{word} → {lemmatizer.lemmatize(word, 'v')}")
```

Output:

running → run

runs → run

ran → run (understands past tense)

easily → easily (preserves adverbs)

fairly → fairly

Differences:

- **Stemming:** Rule-based, faster, sometimes produces non-words ("easili")
- **Lemmatization:** Dictionary-based, accurate, requires POS tags

Modern approach: Skip this for LLMs. Transformers learn representations that handle these naturally.

Word Embeddings: Representing Meaning as Vectors

Embeddings convert discrete tokens into continuous vectors that capture semantic meaning.

Word2Vec: Skip-gram and CBOW

Word2Vec uses shallow neural networks to learn word representations by predicting context words.

Skip-gram model concept

Given: "The quick brown fox"

Predict surrounding words from each word

"quick" → predict ["The", "brown"]
(context window = 2)

"brown" → predict ["quick", "fox"]

Implementation

```
from gensim.models import Word2Vec  
  
sentences = [  
    ["the", "quick", "brown", "fox"],  
    ["the", "lazy", "dog"],  
]  
  
model = Word2Vec(  
    sentences=sentences,  
    vector_size=100, # embedding dimension  
    window=2, # context window size  
    min_count=1, # minimum word frequency  
    sg=1 # 1=Skip-gram, 0=CBOW  
)
```

Access embeddings

```
vector = model.wv["quick"] # shape: (100,)  
similarity = model.wv.similarity("quick", "fast") # cosine similarity
```

Key metrics:

- **Vector dimension:** Typically 100-300 for Word2Vec. Higher = more capacity but more data needed
- **Context window:** 2-5 words on each side. Larger window captures broader context but loses syntactic info
- **Training method:** Skip-gram learns better representations but is slower; CBOW is faster

GloVe: Global Vectors for Word Representation

GloVe combines global matrix factorization with local context windows:

GloVe doesn't require separate tokenization; it builds a co-occurrence matrix

and factorizes it to learn embeddings

Using pre-trained GloVe embeddings

```
import numpy as np
```

Load pre-trained GloVe (e.g., 100-dim trained on 6B tokens)

```
embeddings_index = {}
with open('glove.6B.100d.txt') as f:
    for line in f:
        values = line.split()
        word = values[0]
        coef = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coef
```

Access embedding

```
queen_embedding = embeddings_index['queen'] # shape: (100,
```

Interesting property: word analogies

king - man + woman \approx queen

```
king = embeddings_index['king']
man = embeddings_index['man']
woman = embeddings_index['woman']
queen_predicted = king - man + woman
```

Cosine similarity with actual queen is very high!

Comparison with Word2Vec:

- **GloVe**: Captures both local context and global corpus statistics. Better for similarity tasks
- **Word2Vec**: Simpler, trains faster. Better for specific downstream tasks

Why LLMs don't use static embeddings: Modern Transformer models learn context-dependent embeddings. "bank" in "river bank" vs "savings bank" gets different representations based on context.

Deep Learning Basics: Foundation for Language Models

Neural Networks and Backpropagation

A neural network is a composition of functions that learns from data through gradient descent.

Simplified neural network forward pass

```
import numpy as np

class SimpleNeuralNetwork:
    def __init__(self, input_dim, hidden_dim, output_dim):
        # Initialize weights with small random values
        self.W1 = np.random.randn(input_dim, hidden_dim) * 0.01
        self.b1 = np.zeros((1, hidden_dim))
        self.W2 = np.random.randn(hidden_dim, output_dim) * 0.01
        self.b2 = np.zeros((1, output_dim))

    def relu(self, x):
        """Activation function: ReLU(x) = max(0, x)"""
        return np.maximum(0, x)

    def softmax(self, x):
        """Converts logits to probabilities"""
        exp_x = np.exp(x - np.max(x, axis=1, keepdims=True)) # numerical stability
```

```

        return exp_x / np.sum(exp_x, axis=1, keepdims=True)

def forward(self, X):
    """Forward pass: input → hidden layer → output"""
    self.z1 = np.dot(X, self.W1) + self.b1      # Linear transformation
    self.a1 = self.relu(self.z1)                # Non-linearity
    self.z2 = np.dot(self.a1, self.W2) + self.b2 # Output layer
    self.a2 = self.softmax(self.z2)            # Probabilities
    return self.a2

def backward(self, X, y, learning_rate=0.01):
    """Backpropagation: compute gradients and update weights"""
    batch_size = X.shape[0]

    # Output layer gradient
    dz2 = self.a2 - y # derivative of cross-entropy loss + softmax
    dW2 = np.dot(self.a1.T, dz2) / batch_size
    db2 = np.sum(dz2, axis=0, keepdims=True) / batch_size

    # Hidden layer gradient
    da1 = np.dot(dz2, self.W2.T)
    dz1 = da1 * (self.a1 > 0) # derivative of ReLU
    dW1 = np.dot(X.T, dz1) / batch_size
    db1 = np.sum(dz1, axis=0, keepdims=True) / batch_size

    # Update weights (gradient descent)
    self.W1 -= learning_rate * dW1
    self.b1 -= learning_rate * db1
    self.W2 -= learning_rate * dW2
    self.b2 -= learning_rate * db2

```

Training loop

```

nn = SimpleNeuralNetwork(input_dim=10, hidden_dim=64, output_dim=5)
for epoch in range(100):
    logits = nn.forward(X_train)
    nn.backward(X_train, y_train, learning_rate=0.01)
    if epoch % 20 == 0:

```

```
loss = -np.mean(np.log(logits[np.arange(len(y_train)), y_train.argmax(axis=1)]))  
print(f"Epoch {epoch}, Loss: {loss:.4f}")
```

Key concepts:

- **Forward pass:** Data flows through layers, computing predictions
- **Backpropagation:** Compute gradients by applying chain rule from output to input
- **Gradient descent:** Update weights in direction opposite to gradient
- **Learning rate:** Controls step size. Too high = divergence; too low = slow training

Optimization Algorithms: SGD, Adam, and Beyond

Different optimizers have different convergence properties:

```
import torch  
import torch.optim as optim  
  
model = YourModel()
```

1. Stochastic Gradient Descent (SGD)

```
optimizer_sgd = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
```

Momentum: accumulate past gradients to smooth updates

Good for: simple, stable training

Bad for: can get stuck in local minima, slow convergence

2. Adam (Adaptive Moment Estimation) - Modern standard

```
optimizer_adam = optim.Adam(  
    model.parameters(),  
    lr=1e-3, # typical learning rate  
    betas=(0.9, 0.999), # exponential moving average coefficients  
    eps=1e-8, # numerical stability  
    weight_decay=1e-5 # L2 regularization  
)
```

Maintains per-parameter learning rates

Good for: most modern tasks, handles sparse gradients

Bad for: can be memory-heavy for very large models

3. AdamW (Adam with decoupled weight decay)

```
optimizer_adamw = optim.AdamW(  
    model.parameters(),  
    lr=1e-3,  
    weight_decay=0.01 # decoupled from gradient-based updates  
)
```

Improvement over Adam: better generalization

Used in: GPT models, BERT fine-tuning

Training loop with optimizer

```
for epoch in range(num_epochs):  
    for batch in dataloader:  
        # Forward pass  
        outputs = model(batch)  
        loss = criterion(outputs, targets)  
  
        # Backward pass  
        optimizer.zero_grad() # Clear old gradients  
        loss.backward() # Compute new gradients  
  
        # Gradient clipping (prevents exploding gradients)  
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
```

```
# Update weights
optimizer.step()
```

Comparison:

Algorithm	Convergence	Memory	Generalization	Use Case
SGD	Slow	Low	Good	Simple tasks
Momentum	Faster	Low	Good	Traditional ML
Adam	Fast	Medium	OK	Most modern tasks
AdamW	Fast	Medium	Better	LLM training (preferred)

Loss Functions for Language Modeling

Cross-entropy is the standard loss for predicting probability distributions:

```
import torch.nn.functional as F
```

Language modeling task: predict next token

```
logits = model(input_ids) # shape: (batch_size, seq_len, vocab_size)
targets = next_token_ids # shape: (batch_size, seq_len)
```

Cross-entropy loss: measures divergence between predicted and true distribution

```
loss = F.cross_entropy(
    logits.view(-1, vocab_size), # flatten: (batch_size * seq_len, vocab_size)
    targets.view(-1) # flatten: (batch_size * seq_len)
)
```

What's happening mathematically:

For each position, softmax(logits) gives probability distribution

Cross-entropy = $-\log(P(\text{correct_token}))$

This encourages high probability for the correct token and low for others

Perplexity = $\exp(\text{loss})$

Measures how "surprised" the model is on test data

Lower is better. Perplexity of 10 = model thinks test data is ~10x more likely

```
perplexity = torch.exp(loss)
```

Sequential Models: RNNs, LSTMs, and GRUs

Before Transformers, sequential models processed text one token at a time with hidden state.

Recurrent Neural Networks (RNNs)

RNNs maintain a hidden state that gets updated at each timestep:

```
import torch
import torch.nn as nn

class SimpleRNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super().__init__()
        # Weight matrices
        self.W_h = nn.Linear(hidden_size, hidden_size) # hidden to hidden
        self.W_x = nn.Linear(input_size, hidden_size) # input to hidden
        self.W_o = nn.Linear(hidden_size, output_size) # hidden to output
```

```

def forward(self, x):
    # x shape: (batch_size, seq_len, input_size)
    batch_size, seq_len, _ = x.shape
    hidden = torch.zeros(batch_size, self.W_h.out_features)
    outputs = []

    for t in range(seq_len):
        #  $h_t = \tanh(W_x * x_t + W_h * h_{t-1})$ 
        hidden = torch.tanh(
            self.W_x(x[:, t, :]) + self.W_h(hidden)
        )
        output = self.W_o(hidden)
        outputs.append(output.unsqueeze(1))

    # outputs shape: (batch_size, seq_len, output_size)
    return torch.cat(outputs, dim=1)

```

PyTorch's built-in RNN (more efficient)

```

rnn = nn.RNN(
    input_size=100,
    hidden_size=256,
    num_layers=2, # stack 2 RNN layers
    batch_first=True # input: (batch, seq, features)
)

outputs, hidden = rnn(x) # outputs: (batch, seq, hidden_size)

```

Problem with RNNs: Vanishing/Exploding Gradients

When backpropagating through many timesteps, gradients get multiplied repeatedly:

- Gradient $\approx \partial L / \partial h_0 = (\partial h_1 / \partial h_0) \times (\partial h_2 / \partial h_1) \times \dots \times (\partial h_t / \partial h_{t-1})$
- If each factor < 1 , product $\rightarrow 0$ (vanishing)
- If each factor > 1 , product $\rightarrow \infty$ (exploding)
- Result: Can't learn long-range dependencies

Long Short-Term Memory (LSTM)

LSTMs solve the vanishing gradient problem with gating mechanisms:

```

class LSTMCell(nn.Module):
    """Single LSTM cell"""

    def __init__(self, input_size, hidden_size):
        super().__init__()

```

```

self.input_size = input_size
self.hidden_size = hidden_size

# All gates combined into one computation for efficiency
self.gates = nn.Linear(input_size + hidden_size, 4 * hidden_size)

def forward(self, x, (h, c)):
    # x: (batch, input_size)
    # h: hidden state (batch, hidden_size)
    # c: cell state (batch, hidden_size)

    # Concatenate input and hidden state
    combined = torch.cat([x, h], dim=1)

    # Compute all gates
    gates = self.gates(combined) # (batch, 4*hidden_size)

    # Split into 4 gates
    input_gate, forget_gate, candidate_gate, output_gate = \
        gates.chunk(4, dim=1)

    # Apply sigmoid to gates (values in [0, 1])
    input_gate = torch.sigmoid(input_gate)
    forget_gate = torch.sigmoid(forget_gate)
    output_gate = torch.sigmoid(output_gate)
    candidate_gate = torch.tanh(candidate_gate)

    # Update cell state
    #  $c_t = f_t \odot c_{t-1} + i_t \odot g_t$ 
    c_new = forget_gate * c + input_gate * candidate_gate

    # Compute hidden state
    #  $h_t = o_t \odot \tanh(c_t)$ 
    h_new = output_gate * torch.tanh(c_new)

    return h_new, c_new

```

How LSTMs avoid vanishing gradients:

- **Cell state** acts as a "highway" with additive updates ($\partial c_t / \partial c_{t-1}$ involves addition, not multiplication)
- **Forget gate** controls what information to keep ($f_t \in [0,1]$ usually > 0.5)
- **Gradients** can flow unchanged through cell state: $\partial c_T / \partial c_t$ involves products of values close to 1

PyTorch's built-in LSTM

```

lstm = nn.LSTM(
    input_size=100,
    hidden_size=256,
    num_layers=2,
    batch_first=True,
    dropout=0.2 # apply dropout between layers
)

outputs, (h_n, c_n) = lstm(x)

```

outputs: (batch, seq, hidden_size) - all hidden states

h_n: final hidden state

c_n: final cell state

Gated Recurrent Unit (GRU)

GRUs are a simpler variant of LSTMs with fewer gates:

GRU equations (simplified LSTM)

Reset gate: $r_t = \sigma(W_r [h_{t-1}, x_t])$

Update gate: $z_t = \sigma(W_z [h_{t-1}, x_t])$

Candidate: $\tilde{h}_t = \tanh(W_h [r_t \odot h_{t-1}, x_t])$

Hidden state: $h_t = (1 - z_t) \odot \tilde{h}_t + z_t \odot h_{t-1}$

```
gru = nn.GRU(
    input_size=100,
    hidden_size=256,
    num_layers=2,
    batch_first=True
)

outputs, h_n = gru(x)
```

Comparison:

Metric	RNN	GRU	LSTM
Parameters	Fewest	Medium	Most
Training Speed	Fastest	Faster	Slower
Long Dependencies	Poor	Good	Best
Computational Cost	Low	Medium	High

Why LSTMs/GRUs are less used now:

- Sequential processing is slow (can't parallelize)
- Transformers with attention handle long-range dependencies much better
- Maximum practical sequence length: ~500-1000 tokens (vs 100K+ for Transformers)

2. The Transformer Architecture: Core Concepts

Transformers replaced RNNs by enabling parallel processing of sequences through attention mechanisms.

The Attention Mechanism: Foundation of Everything

Attention allows the model to focus on relevant parts of the input regardless of distance.

Scaled Dot-Product Attention

```
import torch
import torch.nn as nn
import math

def scaled_dot_product_attention(Q, K, V, mask=None):
    """
    Q: Query matrix (batch, seq_len, d_k)
    K: Key matrix (batch, seq_len, d_k)
    V: Value matrix (batch, seq_len, d_v)
```

```

mask: Optional mask for autoregressive (causal) attention
"""
d_k = Q.shape[-1]

# Compute attention scores
# Attention(Q, K, V) = softmax(Q * K^T / sqrt(d_k)) * V
scores = torch.matmul(Q, K.transpose(-2, -1)) / math.sqrt(d_k)
# scores shape: (batch, seq_len, seq_len)

# Apply mask (for causal attention, prevent attending to future tokens)
if mask is not None:
    scores = scores.masked_fill(mask == 0, float('-inf'))

# Normalize attention weights to sum to 1
attention_weights = torch.softmax(scores, dim=-1)

# Apply attention weights to values
output = torch.matmul(attention_weights, V)
# output shape: (batch, seq_len, d_v)

return output, attention_weights

```

Example: Sequence of 4 tokens, embedding dim 64

```

seq_len, d_model = 4, 64
Q = torch.randn(1, seq_len, d_model) # queries
K = torch.randn(1, seq_len, d_model) # keys
V = torch.randn(1, seq_len, d_model) # values

output, weights = scaled_dot_product_attention(Q, K, V)

```

weights shape: (1, 4, 4) - attention from each token to all tokens

Why scale by $\sqrt{d_k}$?

- Dot products can be very large, causing softmax to be dominated by one value
- Scaling prevents saturation and keeps gradients stable
- Variance of $Q \cdot K^T / \sqrt{d_k} \approx 1$ (normalized)

Multi-Head Attention

Instead of computing attention once, compute it multiple times in parallel with different subspaces:

```
class MultiHeadAttention(nn.Module):
    def __init__(self, d_model, num_heads):
        super().__init__()
        assert d_model % num_heads == 0

        self.d_model = d_model
        self.num_heads = num_heads
        self.d_k = d_model // num_heads # dimension of each head

        # Linear projections
        self.W_q = nn.Linear(d_model, d_model)
        self.W_k = nn.Linear(d_model, d_model)
        self.W_v = nn.Linear(d_model, d_model)
        self.W_o = nn.Linear(d_model, d_model)

    def forward(self, Q, K, V, mask=None):
        batch_size = Q.shape[0]
        seq_len = Q.shape[1]

        # Linear projections
        Q = self.W_q(Q) # (batch, seq_len, d_model)
        K = self.W_k(K)
        V = self.W_v(V)

        # Split into multiple heads
        # Reshape: (batch, seq_len, d_model) → (batch, seq_len, num_heads, d_k)
        # Transpose: (batch, num_heads, seq_len, d_k)
        Q = Q.reshape(batch_size, seq_len, self.num_heads, self.d_k).transpose(1, 2)
        K = K.reshape(batch_size, seq_len, self.num_heads, self.d_k).transpose(1, 2)
        V = V.reshape(batch_size, seq_len, self.num_heads, self.d_k).transpose(1, 2)

        # Apply scaled dot-product attention
        attn_output, attn_weights = scaled_dot_product_attention(Q, K, V, mask)
        # attn_output: (batch, num_heads, seq_len, d_k)
```

```

# Concatenate heads
attn_output = attn_output.transpose(1, 2).contiguous()
# (batch, seq_len, num_heads, d_k)
attn_output = attn_output.reshape(batch_size, seq_len, self.d_model)
# (batch, seq_len, d_model)

# Final linear projection
output = self.W_o(attn_output)

return output, attn_weights

```

Usage

```

mha = MultiHeadAttention(d_model=512, num_heads=8)
x = torch.randn(batch_size=32, seq_len=100, d_model=512)
output, weights = mha(x, x, x) # Self-attention: Q=K=V=x

```

Each head operates on $512/8 = 64$ -dimensional subspace

Benefits of multi-head attention:

- Different heads attend to different aspects (syntax, semantics, long-range, local patterns)
- Parallelizable: all heads compute simultaneously
- More expressive than single-head attention

The Transformer Block: Self-Attention + FFN

```

class TransformerBlock(nn.Module):
    def __init__(self, d_model, num_heads, d_ff, dropout=0.1):
        super().__init__()

        # Layer 1: Multi-head self-attention
        self.self_attn = MultiHeadAttention(d_model, num_heads)

        # Layer 2: Feed-forward network (MLP)
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.GELU(), # Smooth ReLU variant
            nn.Linear(d_ff, d_model)

```

```

    )

# Normalization and dropout
self.norm1 = nn.LayerNorm(d_model)
self.norm2 = nn.LayerNorm(d_model)
self.dropout = nn.Dropout(dropout)

def forward(self, x, mask=None):
    # x shape: (batch, seq_len, d_model)

    # Sub-layer 1: Self-attention with residual connection
    attn_output, _ = self.self_attn(x, x, x, mask)
    attn_output = self.dropout(attn_output)
    x = self.norm1(x + attn_output) # Residual + LayerNorm

    # Sub-layer 2: Feed-forward with residual connection
    ffn_output = self.ffn(x)
    ffn_output = self.dropout(ffn_output)
    x = self.norm2(x + ffn_output) # Residual + LayerNorm

return x

```

Stack multiple transformer blocks for deeper model

```

class Transformer(nn.Module):
    def __init__(self, d_model, num_heads, num_layers, d_ff, dropout=0.1):
        super().__init__()
        self.layers = nn.ModuleList([
            TransformerBlock(d_model, num_heads, d_ff, dropout)
            for _ in range(num_layers)
        ])

```

```

def forward(self, x, mask=None):
    for layer in self.layers:
        x = layer(x, mask)
    return x

```

Positional Encoding: Injecting Position Information

Attention doesn't have inherent position awareness. We must add position information explicitly.

Sinusoidal Positional Encoding (Transformers)

```
class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_seq_length=5000):
        super().__init__()

        # Create positional encoding matrix
        pe = torch.zeros(max_seq_length, d_model)
        position = torch.arange(0, max_seq_length).unsqueeze(1).float()

        # Compute dimension indices
        div_term = torch.exp(
            torch.arange(0, d_model, 2).float() *
            -(math.log(10000.0) / d_model)
        )

        # Apply sin to even indices, cos to odd indices
        pe[:, 0::2] = torch.sin(position * div_term)    # even positions
        pe[:, 1::2] = torch.cos(position * div_term)    # odd positions

        # Register as buffer (not trained, but saved with model)
        self.register_buffer('pe', pe.unsqueeze(0))

    def forward(self, x):
        # x shape: (batch, seq_len, d_model)
        x = x + self.pe[:, :x.shape[1], :]
        return x
```

Mathematical formula:

$$PE(pos, 2i) = \sin(pos / 10000^{(2i/d_model)})$$

$$PE(pos, 2i+1) = \cos(pos / 10000^{(2i/d_model)})$$

Properties:

- Each dimension has different frequency
- Creates unique encoding for each position
- Can extrapolate to longer sequences than training

Why sinusoidal encoding?

- Provides absolute position information ($PE_{10} \neq PE_{20}$)
- Encodes relative position information ($PE_{\{pos+k\}}$ can be derived from $PE_{\{pos\}}$)
- Works for any sequence length without retraining
- Hardware-friendly (no learnable parameters)

Rotary Position Embedding (RoPE) - Modern Standard

RoPE is more efficient and better for long sequences:

```
class RotaryPositionEmbedding(nn.Module):
    def __init__(self, d_model):
        super().__init__()
        self.d_model = d_model
        # Precompute inverse frequencies
        inv_freq = 1.0 / (10000 ** (torch.arange(0, d_model, 2).float() / d_model))
        self.register_buffer("inv_freq", inv_freq)

    def forward(self, x, seq_len):
        # x shape: (batch, seq_len, d_model)
        # Compute position indices
        t = torch.arange(seq_len, device=x.device, dtype=self.inv_freq.dtype)

        # Compute frequencies for each position
        freqs = torch.einsum("i,j->ij", t, self.inv_freq)

        # Compute cos and sin
        emb = torch.cat([freqs, freqs], dim=-1)
```

```

cos_emb = emb.cos()[None, :, :]
sin_emb = emb.sin()[None, :, :]

return cos_emb, sin_emb

```

```

def apply_rotary_pos_emb(q, k, cos_emb, sin_emb):
    """Apply rotary embeddings to queries and keys"""
    # Reshape for element-wise operations
    q_rot = torch.stack([-q[..., 1::2], q[..., 0::2]], dim=-1).flatten(-2)
    k_rot = torch.stack([-k[..., 1::2], k[..., 0::2]], dim=-1).flatten(-2)

```

```

q = q * cos_emb + q_rot * sin_emb
k = k * cos_emb + k_rot * sin_emb

return q, k

```

Usage in attention

```

rope = RotaryPositionEmbedding(d_model=512)
cos_emb, sin_emb = rope(x, seq_len)
Q_rotated, K_rotated = apply_rotary_pos_emb(Q, K, cos_emb, sin_emb)

```

Advantages of RoPE:

- Extrapolates better to longer sequences
- Used in modern models: LLaMA, Mistral
- More efficient than sinusoidal encoding

Layer Normalization and Residual Connections

Layer Normalization vs Batch Normalization

LayerNorm: normalize across feature dimension

BatchNorm: normalize across batch dimension

LayerNorm (used in Transformers)

```
x = torch.randn(32, 100, 512) # (batch, seq_len, d_model)
norm = nn.LayerNorm(512) # normalize last dimension
y = norm(x)
```

Each token's 512-dimensional vector is normalized independently

Mean of last dimension = 0, Var = 1

Mathematical:

$$y = \text{gamma} * (x - \text{mean}(x)) / \sqrt{\text{var}(x) + \text{eps}} + \text{beta}$$

gamma and beta are learnable parameters

Why LayerNorm in Transformers?

- 1. Works well with attention (doesn't depend on batch size)**
- 2. Effective for variable-length sequences**
- 3. Stabilizes training of deep models**

Residual Connections (Skip Connections)

Without residual: $y = f(x)$

With residual: $y = x + f(x)$

```
class ResidualBlock(nn.Module):
    def __init__(self, d_model):
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_model * 4)
        self.linear2 = nn.Linear(d_model * 4, d_model)
        self.norm = nn.LayerNorm(d_model)

    def forward(self, x):
        # Pre-normalization (modern, more stable)
        normed = self.norm(x)
        output = self.linear2(F.gelu(self.linear1(normed)))
        return x + output # Residual connection
```

Why residual connections?

1. Gradient flow: $\partial L / \partial x = \partial L / \partial y * \partial y / \partial x = \partial L / \partial y * (1 + \partial f / \partial x)$

Even if $\partial f / \partial x \approx 0$, gradient can still flow through

2. Easier optimization: model learns small corrections to identity function

3. Enables training very deep networks (50+layers)

Full Transformer Architecture: Encoder-Decoder Structure

```
class FullTransformer(nn.Module):
    def __init__(self, vocab_size, d_model, num_heads, num_layers, d_ff, max_seq_len, dropout=0.1):
        super().__init__()

        # Embedding layers
        self.embedding = nn.Embedding(vocab_size, d_model)
```

```
self.positional_encoding = PositionalEncoding(d_model, max_seq_len)

# Encoder (processes input)
self.encoder = Transformer(
    d_model=d_model,
    num_heads=num_heads,
    num_layers=num_layers,
    d_ff=d_ff,
    dropout=dropout
)

# Decoder (generates output, attends to encoder)
self.decoder = Transformer(
    d_model=d_model,
    num_heads=num_heads,
    num_layers=num_layers,
    d_ff=d_ff,
    dropout=dropout
)

# Output projection to vocabulary
self.output_projection = nn.Linear(d_model, vocab_size)

def forward(self, src, tgt, src_mask=None, tgt_mask=None, memory_mask=None):
    # src: source sequence (batch, src_len)
    # tgt: target sequence (batch, tgt_len)

    # Embed and add positional encoding
    src_embed = self.positional_encoding(self.embedding(src))
    tgt_embed = self.positional_encoding(self.embedding(tgt))

    # Encode source
    memory = self.encoder(src_embed, src_mask)

    # Decode with cross-attention to encoder
    # (implementation simplified - actual cross-attention needed)
    tgt_output = self.decoder(tgt_embed, tgt_mask)
```

```

# Project to vocabulary
logits = self.output_projection(tgt_output)

return logits

```

Key Architectural Variants

Model	Type	Attention Pattern	Use Case
BERT	Encoder-only	Bidirectional	Classification, NER, QA
GPT	Decoder-only	Causal (masked)	Text generation
T5	Encoder-decoder	Both types	Seq2seq tasks
RoBERTa	Encoder-only	Bidirectional	General NLP
ELECTRA	Encoder-only	Bidirectional	Efficient pretraining
LLaMA	Decoder-only	Causal + RoPE	Large-scale generation
Mistral	Decoder-only	Causal + RoPE	Fast inference

3. LLM Pre-training and Training Objectives

Pre-training on massive text corpora is what gives LLMs their knowledge and capabilities.

Language Modeling Objectives

Autoregressive (Causal) Language Modeling - GPT

Predict the next token given all previous tokens. This is how GPT models are trained.

```

import torch
import torch.nn as nn

class GPTLanguageModel(nn.Module):
    """Autoregressive language model (like GPT)"""

```

```
def __init__(self, vocab_size, d_model, num_heads, num_layers, d_ff):
    super().__init__()

    self.embedding = nn.Embedding(vocab_size, d_model)
    self.positional_encoding = PositionalEncoding(d_model)
    self.transformer = Transformer(d_model, num_heads, num_layers, d_ff)
    self.lm_head = nn.Linear(d_model, vocab_size)

def forward(self, input_ids, labels=None):
    # input_ids: (batch, seq_len)
    # labels: (batch, seq_len) - target tokens for loss computation

    # Get embeddings with positional encoding
    x = self.embedding(input_ids)
    x = self.positional_encoding(x)

    # Create causal mask (can't attend to future positions)
    seq_len = input_ids.shape[1]
    causal_mask = torch.triu(
        torch.ones(seq_len, seq_len),
        diagonal=1
    ).bool().to(input_ids.device)
    # Matrix of shape (seq_len, seq_len) where:
    # causal_mask[i, j] = True if i < j (can't attend to future)

    # Apply transformer
    x = self.transformer(x, mask=causal_mask)

    # Project to vocabulary logits
    logits = self.lm_head(x) # (batch, seq_len, vocab_size)

    # Compute loss if labels provided
    loss = None
    if labels is not None:
        loss = nn.functional.cross_entropy(
            logits.view(-1, vocab_size),
            labels.view(-1)
```

```
)  
  
return logits, loss
```

Training loop

```
model = GPTLanguageModel(vocab_size=50000, d_model=768,  
num_heads=12, num_layers=12, d_ff=3072)  
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-4)  
  
for epoch in range(num_epochs):  
    for batch in dataloader:  
        input_ids = batch['input_ids'] # (batch, seq_len)  
        labels = batch['labels'] # same as input_ids shifted right  
  
        # Forward pass  
        logits, loss = model(input_ids, labels=labels)  
  
        # Backward pass  
        optimizer.zero_grad()  
        loss.backward()  
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)  
        optimizer.step()
```

How labels are created:

If **input_ids** = [101, 2054, 2003, 1998]

labels = [2054, 2003, 1998, 102] (shifted by 1)

Model learns: given [101], predict 2054

given [101, 2054], predict 2003

etc.

Key properties of causal language modeling:

- Only attend to past and current tokens (no future information)
- Predicts next token given context
- Natural for text generation
- Every position contributes to loss

Masked Language Modeling - BERT

Randomly mask tokens and predict them from context. This is how BERT is trained.

```
class BERTModel(nn.Module):
    """Masked Language Model (like BERT)"""

    def __init__(self, vocab_size, d_model, num_heads, num_layers, d_ff):
        super().__init__()

        self.embedding = nn.Embedding(vocab_size, d_model)
        self.positional_encoding = PositionalEncoding(d_model)
        self.transformer = Transformer(d_model, num_heads, num_layers, d_ff)
        self.mlm_head = nn.Linear(d_model, vocab_size)
        self.mask_token_id = 103 # [MASK] token id

    def create_mlm_labels(self, input_ids, mask_prob=0.15):
        """Create MLM task: randomly mask tokens"""
        labels = input_ids.clone()

        # Identify positions to mask
        random_mask = torch.rand(input_ids.shape) < mask_prob
        special_tokens = (input_ids == 101) | (input_ids == 102) # [CLS], [SEP]
        mask_positions = random_mask & ~special_tokens

        # For masked positions:
        # 80% → replace with [MASK]
        # 10% → replace with random token
        # 10% → keep original (noise)

        masked_input = input_ids.clone()
```

```

for i in range(input_ids.shape[0]):
    positions = torch.where(mask_positions[i])[0]
    for pos in positions:
        rand = torch.rand(1).item()
        if rand < 0.8:
            masked_input[i, pos] = self.mask_token_id
        elif rand < 0.9:
            masked_input[i, pos] = torch.randint(0, 50000, (1,)).item()
        # else: keep original

return masked_input, labels

def forward(self, input_ids, masked_input=None):
    if masked_input is None:
        masked_input = input_ids

    # Get embeddings with positional encoding
    x = self.embedding(masked_input)
    x = self.positional_encoding(x)

    # No causal mask - BERT attends bidirectionally
    x = self.transformer(x, mask=None)

    # Project to vocabulary for each position
    logits = self.mlm_head(x) # (batch, seq_len, vocab_size)

return logits

```

Training

```

model = BERTModel(vocab_size=50000, d_model=768, num_heads=12, num_layers=12,
d_ff=3072)
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-4)

for epoch in range(num_epochs):
    for batch in dataloader:
        input_ids = batch['input_ids']

```

```

# Create MLM task
masked_input, labels = model.create_mlm_labels(input_ids)

# Forward pass
logits = model(input_ids, masked_input=masked_input)

# Compute loss only on masked positions
loss = 0
for i in range(len(labels)):
    mask_positions = torch.where((masked_input[i] == 103))[0]
    if len(mask_positions) > 0:
        pred = logits[i, mask_positions]
        true = labels[i, mask_positions]
        loss += nn.functional.cross_entropy(pred, true)

# Backward pass
optimizer.zero_grad()
loss.backward()
optimizer.step()

```

Comparison:

Aspect	Autoregressive (GPT)	Masked (BERT)
Training objective	Predict next token	Predict masked tokens
Context	Left-to-right (causal)	Bidirectional
Generation	Natural (token-by-token)	Requires special handling
Pretraining speed	Slower (sequential)	Faster
Fine-tuning	Direct for generation	Needs adaptation for generation
Example models	GPT, GPT-2, GPT-3	BERT, RoBERTa

Tokenization Strategies: Byte-Pair Encoding (BPE)

```
BPE reduces vocabulary size while maintaining expressiveness through subword units.  
from tokenizers import Tokenizer, models, normalizers, pre_tokenizers, decoders, trainers
```

1. Create a BPE tokenizer

```
tokenizer = Tokenizer(models.BPE())
```

2. Set preprocessing

```
tokenizernormalizer = normalizers.Sequence([  
    normalizers.NFC(),  
])  
  
tokenizer.pre_tokenizer = pre_tokenizers.ByteLevel(add_prefix_space=True)
```

3. Train on corpus

```
trainer = trainers.BpeTrainer(  
    vocab_size=50000,  
    min_frequency=2,  
    special_tokens=["<|endoftext|>", "<|padding|>"]  
)
```

files: list of training files

```
tokenizer.train(files, trainer)
```

4. Use the tokenizer

```
text = "Hello, how are you?"  
encoded = tokenizer.encode(text)  
print(encoded.tokens) # ['Hello', ',', ' how', ' are', ' you', '?']  
print(encoded.ids) # [256, 11, 512, 890, 1024, 34] (example ids)
```

5. Decode back to text

```
decoded = tokenizer.decode(encoded.ids)  
print(decoded) # "Hello, how are you?"
```

BPE algorithm (simplified):

Start with character-level vocabulary

Repeat until vocab_size reached:

- 1. Find most frequent pair of adjacent tokens**
- 2. Merge them into single token**
- 3. Add to vocabulary**

Example:

Iteration 0: [h, e, l, l, o] (vocab: 256 characters)

**Iteration 1: [h, e, l, l, o] → [h, e, "ll", o]
(merge "l" + "l")**

**Iteration 2: → [h, e, "ll", o] → ["he", "ll", o]
(merge "h" + "e")**

...continues until vocab_size=50000

Why BPE is better than word-level tokenization:

- **Vocabulary size:** 50K (BPE) vs 100K+ (word-level)
- **OOV handling:** Rare words → subword pieces (not "UNK")
- **Efficiency:** Shorter sequences, faster processing
- **Compression:** Common patterns are merged early

Modern alternatives:

WordPiece (used in BERT)

- Similar to BPE but uses likelihood-based merging
- Slightly better for classification tasks

SentencePiece (used in T5, mBART)

- Language-agnostic
- Treats special characters as separate tokens
- Better for multilingual models

```
from sentencepiece import SentencePieceProcessor
sp = SentencePieceProcessor()
sp.Load("model.model")
tokens = sp.EncodeAsIds("Hello world")
```

Large-Scale Distributed Training

Training LLMs requires distributed training across multiple GPUs/TPUs.

```
import torch.distributed as dist
from torch.nn.parallel import DistributedDataParallel as DDP
```

Initialize distributed training

```
dist.init_process_group("nccl") # NVIDIA Collective Communications Library
rank = dist.get_rank()
world_size = dist.get_world_size()
device = torch.device(f"cuda:{rank}")
```

Model on GPU

```
model = GPTLanguageModel(...).to(device)
model = DDP(model, device_ids=[rank])
```

Data loader with DistributedSampler

```
sampler = DistributedSampler(
    dataset,
    num_replicas=world_size,
    rank=rank,
    shuffle=True
)
dataloader = DataLoader(dataset, sampler=sampler, batch_size=32)
```

Training loop (same as single-GPU)

```
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-4)

for epoch in range(num_epochs):
    for batch in dataloader:
        # Forward pass
        logits, loss = model(batch['input_ids'], labels=batch['labels'])

        # Backward pass - gradients are accumulated across all GPUs
        optimizer.zero_grad()
        loss.backward()

        # All-reduce: average gradients across GPUs
        # (DDP does this automatically)

        optimizer.step()
```

Cleanup

```
dist.destroy_process_group()
```

Launching with multiple GPUs:

`torchrun --nproc_per_node=4 train.py`

**Launches 4 processes, one per GPU,
automatically handles synchronization**

Key distributed training concepts:

- **Data parallelism:** Each GPU gets different batch, same model
- **Model parallelism:** Model sharded across GPUs (for very large models)
- **All-reduce:** Synchronize gradients across GPUs (collective operation)
- **Pipeline parallelism:** Different layers on different GPUs

4. Adaptation and Fine-Tuning

Fine-tuning adapts pre-trained models to specific tasks. Full parameter fine-tuning is resource-intensive; parameter-efficient methods enable fine-tuning on consumer hardware.

Full Fine-Tuning: Update All Parameters

Load pre-trained model

```
model = GPTLanguageModel.from_pretrained("gpt2-medium")
```

Unfreeze all parameters

```
for param in model.parameters():
    param.requires_grad = True
```

Fine-tune on task-specific data

```
optimizer = torch.optim.AdamW(model.parameters(), lr=5e-5)

for epoch in range(num_epochs):
    for batch in task_dataloader:
        input_ids = batch['input_ids']
        labels = batch['labels']
```

```
        logits, loss = model(input_ids, labels=labels)
```

```
        optimizer.zero_grad()
```

```
loss.backward()  
torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)  
optimizer.step()
```

Save fine-tuned model

```
model.save_pretrained("./finetuned-model")
```

Requirements:

- Memory: 3x model size (weights + gradients + optimizer states)
- For GPT-2 (1.5B parameters): ~18 GB VRAM needed
- Training time: proportional to learning rate cycles needed

Advantages:

- Highest accuracy (all parameters adapted)
- Simple to implement

Disadvantages:

- Catastrophic forgetting: model loses general knowledge
- Requires large amounts of task data and compute
- Hyperparameter tuning critical (learning rate, warmup, etc.)

LoRA: Low-Rank Adaptation [2023]

Instead of updating all parameters, add small trainable matrices alongside frozen weights.

```
class LoRALayer(nn.Module):  
    """Linear layer with LoRA adaptation"""
```

```
def __init__(self, original_layer, rank=8, lora_alpha=16):  
    super().__init__()  
    self.original_layer = original_layer  
    self.rank = rank  
    self.lora_alpha = lora_alpha  
  
    # Freeze original weights  
    original_layer.weight.requires_grad = False  
    original_layer.bias.requires_grad = False  
  
    # LoRA matrices: decompose update as low-rank  
    #  $\Delta W = (A @ B)$  where  $A \in \mathbb{R}^{(in \times r)}$ ,  $B \in \mathbb{R}^{(r \times out)}$   
    in_features = original_layer.in_features
```

```

out_features = original_layer.out_features

self.lora_a = nn.Parameter(
    torch.randn(in_features, rank) * (2 / (in_features + rank))
)
self.lora_b = nn.Parameter(torch.zeros(rank, out_features))

# Scaling factor
self.scaling = lora_alpha / rank

def forward(self, x):
    # W_eff = W_0 + α/r * (A @ B)
    # Where W_0 is frozen original weight
    lora_out = (x @ self.lora_a) @ self.lora_b * self.scaling
    return self.original_layer(x) + lora_out

```

Inject LoRA into model

```

def add_lora_to_model(model, rank=8, lora_alpha=16):
for name, module in model.named_modules():
if isinstance(module, nn.Linear):
setattr(
model.get_submodule(name.rsplit('.', 1)[0]),
name.rsplit('.', 1)[1],
LoRALayer(module, rank=rank, lora_alpha=lora_alpha)
)

model = GPTLanguageModel.from_pretrained("gpt2-medium")
add_lora_to_model(model, rank=8)

```

Count trainable parameters

```

total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Total: {total_params/1e6:.1f}M, Trainable: {trainable_params/1e6:.1f}M")

```

GPT-2 Medium: 345M total, 0.3M trainable (0.09%)

Training is identical

```
optimizer = torch.optim.AdamW(  
    filter(lambda p: p.requires_grad, model.parameters()),  
    lr=1e-4  
)  
  
for epoch in range(num_epochs):  
    for batch in dataloader:  
        logits, loss = model(batch['input_ids'], labels=batch['labels'])  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()
```

Save only LoRA weights (~1-2 MB)

```
torch.save(  
    {name: param for name, param in model.named_parameters() if param.requires_grad},  
    "lora_weights.pt"  
)
```

Load LoRA weights

```
lora_weights = torch.load("lora_weights.pt")  
for name, param in lora_weights.items():  
    model.get_parameter(name).data.copy_(param)
```

Mathematical foundation:

Full update matrix (345M parameters):

ΔW = [massive matrix]

LoRA decomposition (0.3M parameters):

$\Delta W = A @ B$ where $A \in \mathbb{R}^{(4096 \times 8)}$, $B \in \mathbb{R}^{(8 \times 4096)}$

Memory savings: $(4096 \times 8 + 8 \times 4096) / (4096 \times 4096) \approx 0.1\%$

Computation: $x @ \Delta W$ becomes $(x @ A) @ B$ (still efficient)

LoRA Properties:

- **Rank (r):** Controls expressiveness. r=8 for small tasks, r=64 for complex tasks
- **Alpha:** Scaling factor. Typical ratio alpha/rank = 16/8 = 2
- **Dropout:** Regularization applied to LoRA matrices (typically 0.05)
- **Target modules:** Usually QKV projections in attention, sometimes all linear layers

QLoRA: Quantized LoRA [2023]

Combine 4-bit quantization with LoRA for even more memory efficiency.

```
from bitsandbytes.nn import Linear4bit
```

```
def quantize_and_lora(model, rank=8):
    """Convert model to 4-bit with LoRA"""


```

```
for name, module in model.named_modules():
    if isinstance(module, nn.Linear):
        # Replace with 4-bit quantized version
        new_module = Linear4bit(
            module.in_features,
            module.out_features,
            bias=module.bias is not None,
            compute_dtype=torch.bfloat16, # high precision for computation
            weight_dtype=torch.uint8,    # 4-bit weights
        )

        # Add LoRA adapters on top
        new_module = LoRALayer(new_module, rank=rank)
        setattr(model, name, new_module)
```

Using HuggingFace's built-in QLoRA

```
from peft import LoraConfig, get_peft_model
from transformers import AutoModelForCausalLM
```

Load model in 4-bit

```
model = AutoModelForCausalLM.from_pretrained(
    "meta-llama/Llama-2-7b-hf",
    load_in_4bit=True,
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True,
    device_map="auto",
)
```

Apply LoRA

```
lora_config = LoraConfig(  
    r=16,  
    lora_alpha=32,  
    target_modules=["q_proj", "v_proj"], # only QV layers  
    lora_dropout=0.05,  
    bias="none",  
    task_type="CAUSAL_LM"  
)  
  
model = get_peft_model(model, lora_config)
```

Now fine-tune

```
trainer = transformers.Trainer(  
    model=model,  
    args=transformers.TrainingArguments(  
        output_dir=".output",  
        learning_rate=2e-4,  
        num_train_epochs=3,  
        per_device_train_batch_size=4,  
        gradient_accumulation_steps=4,  
    ),  
    train_dataset=dataset,  
)  
  
trainer.train()
```

Memory comparison:

Model: LLaMA-65B

Full Fine-tuning:

- Model: 130 GB (float16)
- Gradients: 130 GB
- Optimizer states: 130 GB
- Total: ~390 GB (costs \$50,000+ in H100 GPUs)

LoRA (rank=8):

- Model: 130 GB (frozen)
- LoRA weights: ~0.3 MB (trainable)
- Gradients: ~0.3 MB
- Optimizer states: ~0.6 MB
- Total: ~130 GB (needs one A100 GPU: ~\$10,000)

QLoRA:

- Model: 16 GB (4-bit quantized)

- LoRA weights: ~0.3 MB
- Total: ~16 GB (works on RTX 4090: ~\$2000)

QLoRA innovations:

- **4-bit NormalFloat (NF4):** Custom quantization format preserving accuracy
- **Double quantization:** Quantize the quantization constants
- **Paged optimizers:** Swap optimizer states to CPU RAM
- **Result:** Fine-tune 65B parameter models on single RTX 4090

Other Parameter-Efficient Methods

Prefix Tuning

```
class PrefixTuning(nn.Module):
    """Prepend learnable prefix to attention keys/values"""

    def __init__(self, d_model, prefix_len=20, num_heads=12):
        super().__init__()
        self.prefix_len = prefix_len
        self.num_heads = num_heads

        # Learnable prefix (hidden dimension × prefix length)
        self.prefix_embeddings = nn.Parameter(
            torch.randn(prefix_len, d_model)
        )

        # Optional projection network
        self.projection = nn.Sequential(
            nn.Linear(d_model, d_model),
            nn.Tanh(),
            nn.Linear(d_model, num_heads * prefix_len)
        )

    def forward(self, key, value):
        # Prepend prefix to key and value sequences
        prefix = self.prefix_embeddings.unsqueeze(0) # (1, prefix_len, d_model)
        key = torch.cat([prefix, key], dim=1)      # prepend
        value = torch.cat([prefix, value], dim=1)
        return key, value
```

Training: only train prefix_embeddings (~2-5 MB for LLMs)

Comparison of PEFT methods:

Method	Memory	Speed	Accuracy	Flexibility
Full FT	3x	Slow	Best	Full
LoRA	1x	Fast	95-99%	High
QLoRA	0.3x	Fast	95-98%	High
Prefix Tuning	1.1x	Fast	90-95%	Medium
Adapter Tuning	1.05x	Fast	92-96%	Medium
BitFit	1x	Fast	85-90%	Low

Instruction Tuning: Supervised Fine-Tuning (SFT)

Convert pre-trained models to instruction-following models through supervised fine-tuning.

Instruction tuning dataset format

```
dataset = [
{
    "instruction": "Summarize this text in one sentence:",
    "input": "Transformers have revolutionized NLP...",
    "output": "Transformers are the foundation of modern NLP."
},
{
    "instruction": "Translate English to French",
    "input": "Hello, how are you?",
    "output": "Bonjour, comment allez-vous?"
},
# ... thousands more examples
]

class InstructionTuningDataset:
    def __init__(self, examples, tokenizer, max_length=512):
        self.examples = examples
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __getitem__(self, idx):
        example = self.examples[idx]
```

```

# Format: instruction + input + output
prompt = f"{example['instruction']}\n{example['input']}\n"
full_text = prompt + example['output']

# Tokenize
inputs = self.tokenizer(
    full_text,
    max_length=self.max_length,
    truncation=True,
    padding="max_length"
)

# Mask loss for instruction tokens (only learn to predict output)
prompt_tokens = self.tokenizer(
    prompt,
    max_length=self.max_length,
    truncation=True
)
prompt_len = len(prompt_tokens['input_ids'])

# Labels: -100 (ignored by loss) for prompt, token ids for output
labels = [-100] * prompt_len + inputs['input_ids'][prompt_len:]

return {
    'input_ids': inputs['input_ids'],
    'attention_mask': inputs['attention_mask'],
    'labels': labels
}

```

Training

```

dataloader = DataLoader(
    InstructionTuningDataset(dataset, tokenizer),
    batch_size=32,
    shuffle=True
)

model = GPTLanguageModel.from_pretrained("gpt2-medium")
optimizer = torch.optim.AdamW(model.parameters(), lr=5e-5)

```

```
for epoch in range(3):
    for batch in dataloader:
        input_ids = batch['input_ids'].to(device)
        labels = batch['labels'].to(device)

        logits, loss = model(input_ids, labels=labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

After instruction tuning, model follows instructions:

**model.generate("Translate to French:
Hello")**

→ "Bonjour"

RLHF: Reinforcement Learning from Human Feedback [2022]

RLHF is a 3-step process:

- 1. SFT: Train model on high-quality examples (done above)**
- 2. Reward modeling: Train model to predict human preferences**

3. PPO: Optimize policy using reward signal

Step 2: Train reward model

```
class RewardModel(nn.Module):
    def __init__(self, base_model):
        super().__init__()
        self.model = base_model
        self.reward_head = nn.Linear(768, 1) # predict scalar reward

    def forward(self, input_ids):
        hidden = self.model(input_ids)[0] # get hidden states
        reward = self.reward_head(hidden[:, -1, :]) # use [EOS] token representation
        return reward
```

Dataset: pairs of responses with human preference labels

```
preference_dataset = [
    {
        "prompt": "Summarize this article...",
        "chosen": "This article discusses...", # better response
        "rejected": "The article is about..." # worse response
    },
    # ...
]
```

Training: maximize $\log(\sigma(\text{reward_chosen} - \text{reward_rejected}))$

σ is sigmoid, so we want: $\text{reward_chosen} > \text{reward_rejected}$

```
class PreferenceLoss(nn.Module):
    def forward(self, reward_chosen, reward_rejected):
        return -torch.log(torch.sigmoid(reward_chosen - reward_rejected)).mean()
```

Step 3: PPO training (complex, uses policy gradients)

Simplification:

```
from trl import PPOTrainer  
  
ppo_config = PPOConfig(learning_rate=1e-5)  
ppo_trainer = PPOTrainer(  
    model=sft_model,  
    config=ppo_config,  
    dataset=preference_dataset,  
    reward_model=reward_model,  
)  
  
ppo_trainer.train()
```

Modern alternative: DPO (Direct Preference Optimization) [2023]

DPO simplifies RLHF by directly optimizing policy without separate reward model:

```
class DPOLoss(nn.Module):  
    def __init__(self, beta=0.5):  
        self.beta = beta  
  
    def forward(self, policy_chosen_logps, policy_rejected_logps,  
              reference_chosen_logps, reference_rejected_logps):  
        """  
        policy_chosen_logps: log probabilities of chosen response under policy  
        policy_rejected_logps: log probabilities of rejected response under policy  
        reference_*: same but under reference (pre-SFT) model  
        """  
        # DPO objective  
        chosen_log_probs = policy_chosen_logps - reference_chosen_logps  
        rejected_log_probs = policy_rejected_logps - reference_rejected_logps  
  
        loss = -torch.log(  
            torch.sigmoid(  
                self.beta * (chosen_log_probs - rejected_log_probs)  
            )  
        ).mean()
```

return loss

Training is simpler:

For each preference pair:

- 1. Compute logits from both policy and reference model**
- 2. Apply DPO loss**
- 3. Backprop and update**

Result: Models trained with DPO rival RLHF-trained models

Advantage: No separate reward model, simpler pipeline

5. Prompt Engineering and Prompting Techniques

The way you frame a question dramatically affects LLM output quality.

Zero-Shot Prompting

Ask the model to solve a task without examples:

Zero-shot classification

prompt = """Classify the sentiment of this text as positive, negative, or neutral.

Text: "I absolutely love this product! It exceeded all my expectations."

Sentiment:"""

response = model.generate(prompt, max_tokens=10)

Output: "Positive"

Zero-shot arithmetic

```
prompt = """Solve: 15 × 23 = ?"""
response = model.generate(prompt, max_tokens=10)
```

Output: "345"

Zero-shot question answering

```
prompt = """Question: What is the capital of France?
Answer:"""
response = model.generate(prompt, max_tokens=10)
```

Output: "Paris"

Characteristics:

- No task examples provided
- Model relies on general knowledge from pre-training
- Works well for straightforward tasks
- May fail for complex, specialized tasks

Few-Shot Prompting

Provide a few examples to guide the model:

Few-shot sentiment classification

```
prompt = """Classify sentiment as positive or negative.
```

Text: "I love this product!"

Sentiment: positive

Text: "This is terrible."

Sentiment: negative

Text: "It's okay, nothing special."

Sentiment: """"

```
response = model.generate(prompt, max_tokens=10)
```

Output: "neutral" or "negative"

Few-shot arithmetic

```
prompt = """Solve the following:
```

```
9 × 12 = 108
```

```
7 × 8 = 56
```

```
6 × 13 = ?"""
```

```
response = model.generate(prompt, max_tokens=10)
```

Output: "78"

Few-shot translation

```
prompt = """Translate English to French:
```

```
English: "Hello"
```

```
French: "Bonjour"
```

```
English: "Good morning"
```

```
French: "Bon matin"
```

```
English: "How are you?"
```

```
French: """
```

```
response = model.generate(prompt, max_tokens=10)
```

Output: "Comment allez-vous?"

Properties:

- 2-5 examples typically enough
- In-context learning: model learns from examples without parameter updates
- Improves performance significantly vs zero-shot
- Examples should be diverse and representative

Chain-of-Thought (CoT) Prompting

Explicitly ask the model to show reasoning steps:

Without CoT

prompt = """A store sells apples at \$2 each. Sarah buys 5 apples and gets 10% discount. How much does she pay?"""

response = model.generate(prompt, max_tokens=10)

Output might be wrong due to missing intermediate steps

With CoT

prompt = """Let's solve this step by step.

A store sells apples at \$2 each. Sarah buys 5 apples and gets 10% discount. How much does she pay?

Step 1: Calculate original price

Original price = 5 apples × \$2/apple = \$10

Step 2: Calculate discount amount

Discount = 10% of \$10 = \$1

Step 3: Calculate final price

Final price = \$10 - \$1 = \$9

Answer: Sarah pays \$9"""

Model much more likely to get this right!

Few-shot CoT

prompt = """Solve problems step-by-step.

Q: If there are 3 cars in the parking lot and 2 more arrive, how many are there?

A: Let me think step-by-step:

- Started with 3 cars
 - 2 more arrived
 - $3 + 2 = 5$ cars
- Total: 5 cars

Q: A bakery makes 20 cookies per hour. How many in 8 hours?

A: Let me think step-by-step:

- Makes 20 cookies per hour
- Working for 8 hours

- $20 \times 8 = 160$ cookies
Total: 160 cookies

Q: A book costs \$15. Buy 3 with 25% discount. Total cost?
A: """"

Model benefits from CoT pattern in examples

Why CoT works:

- Forces model to break down complex problems
- Makes errors more visible (wrong intermediate step)
- Activates better reasoning capabilities
- Works especially well for math and logic

Tree-of-Thought (ToT) Prompting

Explore multiple reasoning paths and select the best:

```
class TreeOfThought:
def __init__(self, model, branch_factor=3, depth=3):
    self.model = model
    self.branch_factor = branch_factor
    self.depth = depth

    def generate_thoughts(self, context, num_thoughts=3):
        """Generate multiple next thoughts"""
        prompt = f"""Given the context below, generate {num_thoughts} diverse
```

and plausible next thoughts to solve the problem.

Context: {context}

Thought 1: """

```
thoughts = []
for i in range(num_thoughts):
    response = self.model.generate(prompt, max_tokens=100)
    thoughts.append(response)
    prompt += f"\n{response}\nThought {i+2}: "

return thoughts
```

```
def evaluate_thoughts(self, context, thoughts):
```

```
"""Score each thought for viability"""
prompt = f"""For the following problem-solving context and proposed thou
```

rate each thought's quality on a scale of 1-10.

Context: {context}

"""

```
scores = []
for i, thought in enumerate(thoughts):
    prompt += f"Thought {i+1}: {thought}\nRating: "
    response = self.model.generate(prompt, max_tokens=10)
    try:
        score = float(response.split()[0])
    except:
        score = 5.0 # default if parsing fails
    scores.append(score)
```

return scores

```
def solve(self, problem):
    """Solve problem using tree search"""
    # BFS search through thought tree
    queue = [(problem, 0)] # (context, depth)
    solutions = []

    while queue:
        context, depth = queue.pop(0)

        if depth >= self.depth:
            # Reached max depth, extract answer
            solutions.append(context)
            continue

        # Generate next thoughts
        thoughts = self.generate_thoughts(context, self.branch_factor)

        # Evaluate thoughts
        scores = self.evaluate_thoughts(context, thoughts)

        # Keep top thoughts
```

```

ranked = sorted(zip(thoughts, scores), key=lambda x: x[1], reverse=True)
for thought, score in ranked[:self.branch_factor]:
    new_context = context + "\n" + thought
    queue.append((new_context, depth + 1))

# Return best solution
return max(solutions, key=lambda x: self._score_solution(x))

def _score_solution(self, solution):
    """Score final solution"""
    prompt = f"How good is this solution? Rate 1-10.\n{solution}\nRating: "
    response = self.model.generate(prompt, max_tokens=10)
    try:
        return float(response.split()[0])
    except:
        return 5.0

```

Usage

```

tot = TreeOfThought(model, branch_factor=3, depth=2)
answer = tot.solve("Complex reasoning problem...")

```

When to use:

- Complex problems requiring exploration
- Multiple solution paths exist
- Computational budget allows (expensive due to multiple generations)

6. Retrieval-Augmented Generation and Agents

RAG: Retrieving Knowledge During Generation

```

from sentence_transformers import SentenceTransformer
from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

class SimpleRAG:
    def __init__(self, model_name="all-MiniLM-L6-v2"):
        self.encoder = SentenceTransformer(model_name)
        self.documents = []
        self.embeddings = np.array([])

```

```

def add_documents(self, docs):
    """Index documents"""
    self.documents = docs
    self.embeddings = self.encoder.encode(docs)
    # embeddings shape: (num_docs, embedding_dim)

def retrieve(self, query, top_k=3):
    """Find most relevant documents"""
    query_embedding = self.encoder.encode([query])
    similarities = cosine_similarity(query_embedding, self.embeddings)[0]
    top_indices = np.argsort(similarities)[-top_k:][:-1]
    return [self.documents[i] for i in top_indices]

def generate_with_rag(self, query):
    """RAG pipeline"""
    # 1. Retrieve relevant documents
    context_docs = self.retrieve(query, top_k=3)
    context = "\n".join(context_docs)

    # 2. Augment prompt with retrieved context
    prompt = f"""Answer the question based on the provided context.

Context:
{context}

Question: {query}
Answer:"""

    # 3. Generate response
    response = llm.generate(prompt, max_tokens=200)
    return response

```

Usage

```

rag = SimpleRAG()
rag.add_documents([
    "Paris is the capital of France.",
    "The Eiffel Tower is located in Paris.",
    "France is known for wine and cheese."]

```

```
"Machine learning is a subfield of AI.",  
# ... more documents  
])  
  
answer = rag.generate_with_rag("What is Paris famous for?")
```

Retrieves relevant documents about Paris

Generates answer with grounded context

Advanced RAG with vector databases:

```
from langchain.vectorstores import Chroma  
from langchain.embeddings.openai import OpenAIEmbeddings  
from langchain.document_loaders import PDFLoader  
from langchain.text_splitter import RecursiveCharacterTextSplitter
```

1. Load documents

```
loader = PDFLoader("research_paper.pdf")  
documents = loader.load()
```

2. Split into chunks

```
splitter = RecursiveCharacterTextSplitter(  
    chunk_size=1000,  
    chunk_overlap=200  
)  
chunks = splitter.split_documents(documents)
```

3. Create embeddings and store in vector DB

```
embeddings = OpenAIEmbeddings()  
vector_db = Chroma.from_documents(chunks, embeddings, persist_directory="./db")
```

4. Retrieve and generate

```
from langchain.chains import RetrievalQA  
from langchain.llms import OpenAI  
  
retriever = vector_db.as_retriever(search_kwargs={"k": 3})  
  
qa = RetrievalQA.from_chain_type(  
    llm=OpenAI(temperature=0),  
    chain_type="stuff",
```

```
retriever=retriever
)

answer = qa.run("What are the main contributions of this paper?")
```

LLM Agents: Using Tools

Agents use LLMs to decide which tools to call and how to interpret results:

```
import json
from enum import Enum

class ToolType(Enum):
    CALCULATOR = "calculator"
    WEB_SEARCH = "web_search"
    WIKIPEDIA = "wikipedia"

class Tool:
    def __init__(self, name, description, function):
        self.name = name
        self.description = description
        self.function = function

        def execute(self, *args, **kwargs):
            return self.function(*args, **kwargs)
```

```
def calculator(expression):
    """Safely evaluate math expression"""
    try:
        return eval(expression)
    except:
        return "Invalid expression"

def web_search(query):
    """Search the web for information"""
    # Placeholder - would use real API
    return f"Search results for: {query}"

def wikipedia_lookup(topic):
    """Look up topic on Wikipedia"""
    # Placeholder - would use real API
    return f"Wikipedia article about {topic}"
```

Define tools

```
tools = [
    Tool("calculator", "Perform mathematical calculations", calculator),
    Tool("web_search", "Search the internet for information", web_search),
    Tool("wikipedia", "Look up information on Wikipedia", wikipedia_lookup),
]
```

```
class Agent:  
def __init__(self, llm, tools):  
    self.llm = llm  
    self.tools = {tool.name: tool for tool in tools}  
    self.tool_descriptions = "\n".join([  
        f"- {tool.name}: {tool.description}"  
        for tool in tools  
    ])
```

```
def plan(self, query):  
    """Use LLM to decide which tools to use"""  
    prompt = f"""You are a helpful assistant with access to these tools:  
  
{self.tool_descriptions}"""
```

For the user's question, use the tools available to solve it.
Respond in JSON format with your reasoning and tool calls.

Format: `{"thought": "...", "tool": "tool_name", "input": "..."}`

User question: `{query}`
Your response:"""

```
response = self.llm.generate(prompt, max_tokens=500)  
try:  
    return json.loads(response)  
except:  
    return {"thought": "No tool needed", "tool": None, "input": None}
```

```
def execute(self, query):
```

"""Solve query using tools"""

max_iterations = 5

thought_process = []

```
for iteration in range(max_iterations):
```

1. Plan: which tool to use

plan = self.plan(query)

thought_process.append(plan)

2. Check if done

if plan.get("tool") is None:

return plan.get("thought", "Unable to solve")

```

# 3. Execute tool
tool_name = plan["tool"]
if tool_name not in self.tools:
    query = f"Tool {tool_name} not found. Try another."
    continue

tool = self.tools[tool_name]
result = tool.execute(plan["input"])

# 4. Update query with result
query = f"Tool {tool_name} returned: {result}\n\nContinue solving the orig

return "Max iterations reached"

```

Usage

```
agent = Agent(llm, tools)
answer = agent.execute("What's 15 * 23? And what's the current weather in Paris?")
```

Modern agent frameworks:

LangChain agents

```
from langchain.agents import AgentExecutor, create_react_agent
from langchain.tools import tool
from langchain.prompts import PromptTemplate

@tool
def multiply(a: float, b: float) -> float:
    """Multiply two numbers"""
    return a * b

@tool
def web_search(query: str) -> str:
    """Search the web"""
    return f"Results for {query}"

tools = [multiply, web_search]
```

Create agent

```
from langchain.llms import OpenAI
llm = OpenAI(temperature=0)

agent = create_react_agent(llm, tools)
executor = AgentExecutor.from_agent_and_tools(agent=agent, tools=tools)

result = executor.invoke({"input": "What's 15 * 23?"})
```

7. Evaluation, Optimization, and Safety

Evaluation Metrics

Perplexity: How Surprised is the Model?

```
import torch
import torch.nn.functional as F

def calculate_perplexity(model, dataloader):
    """Perplexity = exp(average_cross_entropy_loss)"""
    total_loss = 0
    total_tokens = 0
```

```
model.eval()
with torch.no_grad():
    for batch in dataloader:
        input_ids = batch['input_ids']
        labels = batch['labels']

        logits = model(input_ids)
        loss = F.cross_entropy(
            logits.view(-1, vocab_size),
            labels.view(-1),
            reduction='sum'
        )

        total_loss += loss.item()
        total_tokens += labels.numel()

    average_loss = total_loss / total_tokens
    perplexity = torch.exp(torch.tensor(average_loss))
```

```
    return perplexity
```

Interpretation:

Perplexity = 10: Model thinks test data is ~10x more surprising than training data

Lower is better

GPT-2: ~20-30 on WikiText

GPT-3: ~10-15 on various benchmarks

BLEU Score: Comparing Translations

```
from nltk.translate.bleu_score import sentence_bleu, corpus_bleu
```

BLEU-4 (most common)

```
reference = [["the", "cat", "is", "on", "the", "mat"]]
candidate = ["the", "cat", "is", "on", "mat"]

bleu = sentence_bleu(reference, candidate, weights=(0.25, 0.25, 0.25, 0.25))
print(f"BLEU score: {bleu:.4f}")
```

Corpus BLEU

```
references = [
    [["the", "cat"], ["a", "cat"]], # Multiple valid translations
    [["dog", "runs"]],
]
candidates = [
    ["the", "cat"],
    ["dog", "runs"],
]
corpus_bleu_score = corpus_bleu(references, candidates)
```

ROUGE: Summary Evaluation

```
from rouge_score import rouge_scorer  
scorer = rouge_scorer.RougeScorer(['rouge1', 'rougeL'], use_stemmer=True)
```

ROUGE-1: overlap of unigrams

ROUGE-L: longest common subsequence

```
reference = "The quick brown fox jumps over the lazy dog"  
prediction = "A fast brown fox jumps over the lazy dog"
```

```
scores = scorer.score(reference, prediction)  
print(f"ROUGE-1 F1: {scores['rouge1'].fmeasure:.4f}")  
print(f"ROUGE-L F1: {scores['rougeL'].fmeasure:.4f}")
```

LLM-as-Judge: Using LLMs to Evaluate LLMs

```
def llm_as_judge(question, model_response, reference_response=None):  
    """Use GPT-4 to evaluate model output"""
```

```
if reference_response:  
    prompt = f"""You are an expert evaluator. Rate the following response on a scale of 1-10.  
    Response: {model_response}  
    Question: {question}  
    Reference (ideal) response: {reference_response}  
    Model response: {model_response}  
    Evaluation criteria:  
    • Correctness: Does it answer the question accurately?  
    • Completeness: Does it cover all important points?  
    • Clarity: Is it well-written and easy to understand?  
    Rating (1-10):"""
```

```
else:  
    prompt = f"""You are an expert evaluator. Rate the following response on a scale of 1-10.  
    Response: {model_response}  
    Question: {question}  
    Evaluation criteria:  
    • Correctness and accuracy  
    • Completeness and relevance  
    • Clarity and presentation
```

Rating (1-10):""""

```
rating_response = llm.generate(prompt, max_tokens=10)
try:
    rating = int(rating_response.split()[0])
    return min(10, max(1, rating)) # Clamp to 1-10
except:
    return 5 # Default if parsing fails
```

Advantages of LLM-as-Judge:

- + Can evaluate on complex criteria humans care about
- + No need for human annotations
- Biased toward generator LLM style
- Can be gamed

Model Optimization

Quantization: Reducing Precision

```
import torch
from torch.quantization import quantize_dynamic
```

Dynamic quantization: quantize weights, keep activations in float32

```
model = GPTLanguageModel.from_pretrained("gpt2-medium")
quantized_model = quantize_dynamic(
    model,
    {torch.nn.Linear}, # which modules to quantize
    dtype=torch.qint8 # quantize to int8
)
```

Results:

Model size: 345 MB → 86 MB (4× reduction)

Speed: Slightly slower (dequantization overhead)

Accuracy: <1% degradation on most tasks

Pruning: Removing Unimportant Parameters

```
import torch.nn.utils.prune as prune  
  
model = GPTLanguageModel.from_pretrained("gpt2-medium")
```

Magnitude pruning: remove smallest weights

```
for name, module in model.named_modules():  
    if isinstance(module, torch.nn.Linear):  
        prune.l1_unstructured(module, name='weight', amount=0.3)  
        # Remove 30% of weights with smallest magnitude  
        prune.remove(module, 'weight') # Make permanent
```

Results:

Model size: ~30% reduction

Speed: ~2× faster (fewer operations)

Accuracy: 2-5% degradation

Fine-tuning after pruning can recover accuracy

Distillation: Training Small Models with Large Models

```
class StudentModel(nn.Module):
    """Small model to distill knowledge into"""
    def __init__(self, vocab_size):
        super().__init__()
        # Smaller config: 12M parameters vs 345M teacher
        self.embedding = nn.Embedding(vocab_size, 256)
        self.transformer = Transformer(
            d_model=256,
            num_heads=8,
            num_layers=4, # vs 12 in teacher
            d_ff=1024
        )
        self.lm_head = nn.Linear(256, vocab_size)
```

```
def forward(self, input_ids):
    x = self.embedding(input_ids)
    x = self.positional_encoding(x)
    x = self.transformer(x)
    logits = self.lm_head(x)
    return logits
```

```
def distillation_loss(student_logits, teacher_logits, temperature=4.0):
```

```
....
```

```
Soft targets from teacher guide student training
temperature: higher = softer probability distribution
....
```

```
# Soft targets from teacher
teacher_probs = F.softmax(teacher_logits / temperature, dim=-1)
```

```
# Student log probabilities
student_log_probs = F.log_softmax(student_logits / temperature, dim=-1)
```

```
# KL divergence
kl_loss = F.kl_div(student_log_probs, teacher_probs, reduction='batchmean')
```

```
return kl_loss * (temperature ** 2)
```

Training loop

```
teacher = GPTLanguageModel.from_pretrained("gpt2-medium")
teachereval() # Don't update teacher

student = StudentModel(vocab_size=50000)
optimizer = torch.optim.AdamW(student.parameters(), lr=1e-4)

for epoch in range(num_epochs):
    for batch in dataloader:
        input_ids = batch['input_ids']
        labels = batch['labels']

        # Forward passes
        with torch.no_grad():
            teacher_logits = teacher(input_ids)

        student_logits = student(input_ids)

        # Combined loss: distillation + task-specific
        alpha = 0.7
        distill_loss = distillation_loss(student_logits, teacher_logits)
        task_loss = F.cross_entropy(student_logits.view(-1, vocab_size),
                                   labels.view(-1))

        loss = alpha * distill_loss + (1 - alpha) * task_loss

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

Results:

Student: 12M vs 345M parameters (28× smaller)

Speed: 10× faster

Accuracy: 85-95% of teacher performance

Safety: Detecting and Mitigating Harmful Outputs

Hallucination Detection

```
class HallucinationDetector:  
    """Detect if model is making up facts"""
```

```
def __init__(self, knowledge_base):  
    self.kb = knowledge_base # Facts we know are true  
  
def check_factuality(self, text):  
    """Extract facts and verify against KB"""  
    # 1. Extract factual claims from text  
    facts = self.extract_facts(text)  
  
    # 2. Verify each fact  
    hallucinations = []  
    for fact in facts:  
        if not self.verify_fact(fact):  
            hallucinations.append(fact)  
  
    return hallucinations  
  
def extract_facts(self, text):  
    """Simple extraction: proper noun + predicate"""  
    # Use NER + dependency parsing  
    # Returns: [(subject, verb, object), ...]  
    pass  
  
def verify_fact(self, fact):  
    """Check if fact exists in KB"""  
    subject, verb, obj = fact  
    return (subject, verb, obj) in self.kb
```

Usage

```
detector = HallucinationDetector(knowledge_base)
response = llm.generate("Who won the 2024 Nobel Prize?")
hallucinations = detector.check_factuality(response)

if hallucinations:
    print(f"Warning: Potential hallucinations detected: {hallucinations}")
```

Jailbreak Prevention

```
class SafetyFilter:  
    """Filter harmful outputs"""
```

```
def __init__(self):
    self.harmful_keywords = [
        "bomb", "illegal", "harmful", "dangerous", ...
    ]
    self.unsafe_intents = [
        "request_illegal_activity",
        "request_violence",
        "request_hate_speech",
    ]
```

```
def is_safe(self, text):
    """Check if text contains harmful content"""

    # 1. Keyword filtering
    text_lower = text.lower()
    for keyword in self.harmful_keywords:
        if keyword in text_lower:
            return False

    # 2. Semantic safety check with classifier
    intent = self.classify_intent(text)
    if intent in self.unsafe_intents:
        return False

    return True
```

```
def classify_intent(self, text):  
    """Use safety classifier"""
```

```
# Could be a small fine-tuned model  
# Returns: safe, potentially_harmful, clearly_harmful  
pass
```

Usage

```
safety_filter = SafetyFilter()  
  
prompt = "How do I make a bomb?"  
if not safety_filter.is_safe(prompt):  
    print("Refusing to process unsafe prompt")  
else:  
    response = llm.generate(prompt)
```

Bias Detection

```
import numpy as np  
  
def measure_gender_bias(model, words_male, words_female, template):  
    """  
    Measure association bias between gender words and target words  
    Example: measure association of gender with STEM fields  
    """  
  
    results = {"male": {}, "female": {}}  
  
    for target_word in target_words:  
        # Prompt with male-associated words  
        prompt_male = template.format(gender_word=random.choice(words_male),  
                                       target=target_word)  
        prob_male = model.get_probability_of_next_word(prompt_male, target_word)  
  
        # Prompt with female-associated words  
        prompt_female = template.format(gender_word=random.choice(words_female),  
                                         target=target_word)  
        prob_female = model.get_probability_of_next_word(prompt_female, target_word)  
  
        results["male"][target_word] = prob_male  
        results["female"][target_word] = prob_female  
  
    # Compute bias metric  
    bias = np.mean([
```

```

    results["male"][w] - results["female"][w]
    for w in target_words
)
return bias # 0 = no bias, >0 = male bias, <0 = female bias

```

Example usage

```

words_male = ["man", "boy", "father", "son"]
words_female = ["woman", "girl", "mother", "daughter"]
target_words = ["engineer", "nurse", "CEO", "teacher"]
template = "The {gender_word} is a {target}"

bias = measure_gender_bias(model, words_male, words_female, template)
print(f"Gender bias: {bias:.4f}")

```

8. Deployment and Serving

Inference Optimization: vLLM

vLLM is a fast inference library that optimizes memory usage and throughput:

```

pip install vllm
from vllm import LLM, SamplingParams

```

Load model

```

llm = LLM(
model="meta-llama/Llama-2-7b-hf",
tensor_parallel_size=2, # distributed across 2 GPUs
dtype="float16",
gpu_memory_utilization=0.9,
)

```

Configure sampling

```

sampling_params = SamplingParams(
temperature=0.7,
top_p=0.95,
max_tokens=256,
)

```

Batch inference

```
prompts = [  
    "What is machine learning?",  
    "Explain deep learning",  
    "What are transformers?",  
]  
  
outputs = llm.generate(prompts, sampling_params)  
  
for output in outputs:  
    print(f"Prompt: {output.prompt}")  
    print(f"Generated: {output.outputs[0].text}\n")
```

Performance improvements with vLLM:

- Paged Attention: GPU memory as virtual memory (4× throughput)
- Continuous batching: dynamic request scheduling
- KV cache optimizations
- Quantization support

API Deployment

```
from fastapi import FastAPI  
from pydantic import BaseModel  
  
app = FastAPI()
```

Load model once at startup

```
llm = LLM(model="meta-llama/Llama-2-7b-hf")  
  
class GenerationRequest(BaseModel):  
    prompt: str  
    max_tokens: int = 256  
    temperature: float = 0.7  
  
class GenerationResponse(BaseModel):  
    prompt: str
```

```
generated_text: str

@app.post("/generate", response_model=GenerationResponse)
async def generate(request: GenerationRequest):
    """Generate text from prompt"""

    sampling_params = SamplingParams(
        temperature=request.temperature,
        max_tokens=request.max_tokens,
    )

    outputs = llm.generate(
        [request.prompt],
        sampling_params,
    )

    return GenerationResponse(
        prompt=request.prompt,
        generated_text=outputs[0].outputs[0].text,
    )
```

Run: uvicorn app:app --host 0.0.0.0 --port 8000

Usage:

```
curl -X POST "http://localhost:8000/generate
" \
-H "Content-Type: application/json" \
```

```
-d '{"prompt": "What is AI?", "max_tokens": 256}'
```

Container Deployment (Docker)

Dockerfile

```
FROM nvidia/cuda:12.1.0-runtime-ubuntu22.04  
WORKDIR /app
```

Install dependencies

```
RUN apt-get update && apt-get install -y python3-pip  
COPY requirements.txt .  
RUN pip install -r requirements.txt
```

Copy application

```
COPY app.py .  
COPY model/ ./model/
```

Download model weights (if not included)

```
RUN python3 -c "from transformers import AutoTokenizer, AutoModelForCausalLM;  
AutoTokenizer.from_pretrained('meta-llama/Llama-2-7b-hf');  
AutoModelForCausalLM.from_pretrained('meta-llama/Llama-2-7b-hf')"
```

Expose port

```
EXPOSE 8000
```

Start server

```
CMD ["uvicorn", "app:app", "--host", "0.0.0.0", "--port", "8000"]
```

Build and run

```
docker build -t my-llm-app .  
docker run --gpus all -p 8000:8000 my-llm-app
```

9. Training Your Model on Custom Data

Complete step-by-step guide to train a language model on your own data.

Step 1: Data Preparation

```
import os
import json
from datasets import Dataset, DatasetDict
from pathlib import Path

class DataPrepipeline:
    def __init__(self, data_path, tokenizer, max_seq_len=512):
        self.data_path = Path(data_path)
        self.tokenizer = tokenizer
        self.max_seq_len = max_seq_len

    def load_raw_data(self):
        """Load raw text files"""
        raw_texts = []

        for file_path in self.data_path.glob("*.txt"):
            with open(file_path, 'r', encoding='utf-8') as f:
                raw_texts.append(f.read())

        return raw_texts

    def clean_text(self, text):
        """Clean and normalize text"""
        # Remove extra whitespace
        text = ' '.join(text.split())

        # Remove control characters
        text = ''.join(ch for ch in text if ord(ch) >= 32 or ch in '\n\t')

        # Remove very long sequences (likely corrupted)
        if len(text) > 1e7:
            text = text[:int(1e7)]

        return text

    def tokenize_function(self, examples):
```

```

"""Tokenize text with special handling for document boundaries"""
# Add document separator token
tokenized = self.tokenizer(
    examples['text'],
    truncation=False, # Don't truncate yet
    return_attention_mask=False,
)

# Concatenate all tokens with </s> separator
all_ids = []
for token_ids in tokenized['input_ids']:
    all_ids.extend(token_ids + [self.tokenizer.eos_token_id])

# Create chunks of max_seq_len
total_len = len(all_ids)
chunk_len = (total_len // self.max_seq_len) * self.max_seq_len
all_ids = all_ids[:chunk_len]

chunks = [all_ids[i:i+self.max_seq_len] for i in range(0, len(all_ids), self.max_seq_len)]

return {
    'input_ids': chunks,
    'labels': chunks.copy(), # For language modeling, labels = input_ids shifted
}
}

def prepare_dataset(self, train_split=0.9, val_split=0.05, test_split=0.05):
    """Prepare train/val/test splits"""

    # 1. Load raw data
    raw_texts = self.load_raw_data()
    print(f"Loaded {len(raw_texts)} files")

    # 2. Clean data
    cleaned_texts = [self.clean_text(text) for text in raw_texts]
    print(f"Cleaned {len(cleaned_texts)} texts")

    # 3. Create dataset
    dataset = Dataset.from_dict({'text': cleaned_texts})

```

```
# 4. Tokenize
tokenized = dataset.map(
    self.tokenize_function,
    batched=True,
    batch_size=1000,
    remove_columns=['text'],
)

# 5. Split
splits = tokenized.train_test_split(
    train_size=train_split,
    test_size=1 - train_split,
    seed=42
)

# Further split test set into val and test
val_test = splits['test'].train_test_split(
    train_size=val_split / (1 - train_split),
    seed=42
)

final_dataset = DatasetDict({
    'train': splits['train'],
    'validation': val_test['train'],
    'test': val_test['test'],
})

print(f"Dataset split:")
print(f" Train: {len(final_dataset['train'])} samples")
print(f" Val: {len(final_dataset['validation'])} samples")
print(f" Test: {len(final_dataset['test'])} samples")

return final_dataset
```

Usage

```
from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("gpt2")
prep = DataPrepipeline("./data/raw", tokenizer, max_seq_len=512)
dataset = prep.prepare_dataset()
```

Step 2: Initialize Model

```
from transformers import AutoConfig, AutoModelForCausalLM
import torch
```

```
class ModelInitializer:
    @staticmethod
    def create_custom_config(model_size="small"):
        """Create custom model configuration"""

        configs = {
```

```
        "small": {
            "hidden_size": 256,
            "num_hidden_layers": 4,
            "num_attention_heads": 8,
            "intermediate_size": 1024,
        },
        "medium": {
            "hidden_size": 512,
            "num_hidden_layers": 8,
            "num_attention_heads": 8,
            "intermediate_size": 2048,
        },
        "large": {
            "hidden_size": 768,
            "num_hidden_layers": 12,
            "num_attention_heads": 12,
            "intermediate_size": 3072,
        },
    }
```

```
config_dict = configs[model_size]
config = AutoConfig.from_pretrained("gpt2")
```

```
for key, value in config_dict.items():
    setattr(config, key, value)

return config

@staticmethod
def initialize_weights(model, std=0.02):
    """Initialize weights (important for training stability)"""
    for name, param in model.named_parameters():
        if "weight" in name and param.dim() > 1:
            torch.nn.init.normal_(param, mean=0.0, std=std)
        elif "bias" in name:
            torch.nn.init.constant_(param, 0.0)

@staticmethod
def create_model(model_size="medium", use_pretrained_weights=False):
    """Create model from scratch or with pretrained weights"""

    if use_pretrained_weights:
        model = AutoModelForCausalLM.from_pretrained("gpt2-medium")
        print("Loaded pretrained GPT-2 weights")
    else:
        config = ModelInitializer.create_custom_config(model_size)
        model = AutoModelForCausalLM.from_config(config)
        ModelInitializer.initialize_weights(model)
        print(f"Created {model_size} model from scratch")

    # Count parameters
    total_params = sum(p.numel() for p in model.parameters())
    trainable_params = sum(p.numel() for p in model.parameters() if p.requires_
        print(f"Model size: {total_params/1e6:.1f}M parameters ({trainable_params/1e6:.1f}M trainable parameters)")

    return model
```

Usage

```
model = ModelInitializer.create_model(model_size="medium",
use_pretrained_weights=False)
```

Step 3: Training Configuration

```
from transformers import TrainingArguments, Trainer
import wandb
```

Initialize weights & biases for monitoring

```
wandb.init(project="llm-training")

training_args = TrainingArguments(
    output_dir="./outputs",

    # Training
    num_train_epochs=3,
    per_device_train_batch_size=32,
    per_device_eval_batch_size=32,
    gradient_accumulation_steps=1,

    # Optimization
    learning_rate=5e-5,
    warmup_steps=1000,
    weight_decay=0.01,
    lr_scheduler_type="cosine",
    max_grad_norm=1.0,

    # Evaluation
    evaluation_strategy="steps",
    eval_steps=500,
    save_strategy="steps",
    save_steps=500,
    save_total_limit=3,

    # Logging
    logging_dir=".//logs",
    logging_steps=100,
```

```
# Hardware
fp16=True, # Mixed precision
gradient_checkpointing=True, # Save memory
dataloader_num_workers=4,

# Misc
seed=42,
report_to="wandb",

)
```

Initialize trainer

```
trainer = Trainer(
model=model,
args=training_args,
train_dataset=dataset['train'],
eval_dataset=dataset['validation'],
data_collator=transformers.default_data_collator,
)
```

Step 4: Training Loop

Train

```
train_result = trainer.train()
```

Evaluate

```
eval_results = trainer.evaluate()
print(f"Validation loss: {eval_results['eval_loss']:.4f}")
print(f"Validation perplexity: {torch.exp(torch.tensor(eval_results['eval_loss'])):.4f}")
```

Save model

```
model.save_pretrained("./final-model")
tokenizers.save_pretrained("./final-model")
```

Test

```
test_results = trainerevaluate(eval_dataset=dataset['test'])
print(f"Test loss: {test_results['eval_loss']:.4f}")
```

Step 5: Inference on Custom Model

```
from transformers import pipeline
```

Load trained model

```
model = AutoModelForCausalLM.from_pretrained("./final-model")
tokenizer = AutoTokenizer.from_pretrained("./final-model")
```

Create text generation pipeline

```
generator = pipeline("text-generation", model=model, tokenizer=tokenizer, device=0)
```

Generate

```
prompt = "The future of artificial intelligence"
output = generator(
    prompt,
    max_length=256,
    num_return_sequences=3,
    temperature=0.7,
    top_p=0.9,
    do_sample=True,
)

for i, seq in enumerate(output):
    print(f"\nGeneration {i+1}:")
    print(seq['generated_text'])
```

10. Replicating GPT-2 from Scratch

Building GPT-2 from first principles to understand architecture deeply.

Architecture Overview

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import math

class GPT2Config:
    """Configuration for GPT-2"""
    def __init__(
```

```
self,
vocab_size=50257,
block_size=1024,
n_layer=12,
n_head=12,
n_embd=768,
dropout=0.1,
):
    self.vocab_size = vocab_size
    self.block_size = block_size # Maximum sequence length
    self.n_layer = n_layer # Number of transformer blocks
    self.n_head = n_head # Number of attention heads
    self.n_embd = n_embd # Embedding dimension
    self.dropout = dropout
```

```
class CausalSelfAttention(nn.Module):
    """Multi-head self-attention with causal mask""""
```

```
def __init__(self, config):
    super().__init__()
    assert config.n_embd % config.n_head == 0

    self.n_head = config.n_head
    self.n_embd = config.n_embd
    self.head_dim = config.n_embd // config.n_head

    # Project to Q, K, V all at once
    self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
    self.c_proj = nn.Linear(config.n_embd, config.n_embd)

    # Dropout
    self.attn_dropout = nn.Dropout(config.dropout)
    self.proj_dropout = nn.Dropout(config.dropout)

    # Register causal mask
    self.register_buffer(
        "bias",
        torch.tril(torch.ones(config.block_size, config.block_size)).view(
            1, 1, config.block_size, config.block_size
        )
    )
```

```

def forward(self, x):
    B, T, C = x.shape # batch_size, seq_len, n_embd

    # Project to Q, K, V
    qkv = self.c_attn(x) # (B, T, 3*C)
    q, k, v = qkv.split(self.n_embd, dim=2) # Each (B, T, C)

    # Reshape for multi-head: (B, T, C) → (B, T, n_head, head_dim)
    q = q.reshape(B, T, self.n_head, self.head_dim).transpose(1, 2)
    k = k.reshape(B, T, self.n_head, self.head_dim).transpose(1, 2)
    v = v.reshape(B, T, self.n_head, self.head_dim).transpose(1, 2)
    # Now: (B, n_head, T, head_dim)

    # Scaled dot-product attention
    scores = torch.matmul(q, k.transpose(-2, -1)) / math.sqrt(self.head_dim)
    # scores: (B, n_head, T, T)

    # Apply causal mask
    scores = scores.masked_fill(self.bias[:, :, :T, :T] == 0, float('-inf'))

    # Softmax
    attn_weights = F.softmax(scores, dim=-1)
    attn_weights = self.attn_dropout(attn_weights)

    # Apply to values
    attn_output = torch.matmul(attn_weights, v) # (B, n_head, T, head_dim)

    # Concatenate heads
    attn_output = attn_output.transpose(1, 2).contiguous()
    # (B, T, n_head, head_dim)
    attn_output = attn_output.reshape(B, T, C) # (B, T, C)

    # Project
    out = self.c_proj(attn_output)
    out = self.proj_dropout(out)

return out

```

```
class MLP(nn.Module):
    """Feed-forward network in transformer block"""

    def __init__(self, config):
        super().__init__()
        self.c_fc = nn.Linear(config.n_embd, 4 * config.n_embd) # Expand
        self.c_proj = nn.Linear(4 * config.n_embd, config.n_embd) # Contract
        self.dropout = nn.Dropout(config.dropout)

    def forward(self, x):
        x = F.gelu(self.c_fc(x)) # GELU activation
        x = self.c_proj(x)
        x = self.dropout(x)
        return x
```

```
class TransformerBlock(nn.Module):
    """Single transformer block"""

    def __init__(self, config):
        super().__init__()
        self.ln_1 = nn.LayerNorm(config.n_embd)
        self.attn = CausalSelfAttention(config)
        self.ln_2 = nn.LayerNorm(config.n_embd)
        self.mlp = MLP(config)

    def forward(self, x):
        # Pre-normalization
        x = x + self.attn(self.ln_1(x)) # Residual connection
        x = x + self.mlp(self.ln_2(x))
        return x
```

```
class GPT2(nn.Module):
    """Full GPT-2 model"""

    def __init__(self, config):
        super().__init__()
        self.config = config
```

```

# Embeddings
self.transformer = nn.ModuleDict(dict(
    wte=nn.Embedding(config.vocab_size, config.n_embd), # Token embedding
    wpe=nn.Embedding(config.block_size, config.n_embd), # Position embedding
    drop=nn.Dropout(config.dropout),
    h=nn.ModuleList([TransformerBlock(config) for _ in range(config.n_layer)])
    ln_f=nn.LayerNorm(config.n_embd),
))

# Language modeling head
self.lm_head = nn.Linear(config.n_embd, config.vocab_size)

# Weight tying: share weights between embedding and output projection
self.transformer.wte.weight = self.lm_head.weight

# Initialize weights
self.apply(self._init_weights)

def _init_weights(self, module):
    if isinstance(module, nn.Linear):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)
        if module.bias is not None:
            torch.nn.init.zeros_(module.bias)
    elif isinstance(module, nn.Embedding):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)

def forward(self, idx, targets=None):
    B, T = idx.shape

    # Get embeddings
    token_emb = self.transformer.wte(idx) # (B, T, n_embd)
    pos_emb = self.transformer.wpe(torch.arange(T, device=idx.device)) # (T, n_embd)
    x = self.transformer.drop(token_emb + pos_emb)

    # Apply transformer blocks
    for block in self.transformer.h:
        x = block(x)

```

```

# Final layer norm
x = self.transformer.ln_f(x)

# Project to vocabulary
logits = self.lm_head(x) # (B, T, vocab_size)

# Compute loss if targets provided
loss = None
if targets is not None:
    loss = F.cross_entropy(
        logits.view(-1, self.config.vocab_size),
        targets.view(-1)
    )

return logits, loss

@torch.no_grad()
def generate(self, idx, max_new_tokens, temperature=1.0, top_k=None):
    """Generate tokens autoregressively"""
    for _ in range(max_new_tokens):
        # Crop input to block_size
        idx_cond = idx if idx.shape[1] <= self.config.block_size else idx[:, -self.config.block_size:]

        # Forward pass
        logits, _ = self(idx_cond)

        # Get logits for next token
        logits = logits[:, -1, :] / temperature

        # Optional top-k sampling
        if top_k is not None:
            v, _ = torch.topk(logits, top_k)
            logits[logits < v[:, [-1]]] = float('-inf')

        # Sample
        probs = F.softmax(logits, dim=-1)
        idx_next = torch.multinomial(probs, num_samples=1)

```

```
# Append  
idx = torch.cat([idx, idx_next], dim=1)  
  
return idx
```

Training

```
config = GPT2Config()  
model = GPT2(config)  
optimizer = torch.optim.AdamW(model.parameters(), lr=6e-4)  
  
for epoch in range(num_epochs):  
    for batch_idx, (X, Y) in enumerate(dataloader):  
        logits, loss = model(X, Y)  
  
        optimizer.zero_grad()  
        loss.backward()  
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)  
        optimizer.step()  
  
        if batch_idx % 100 == 0:  
            print(f"Epoch {epoch}, Batch {batch_idx}, Loss: {loss.item():.4f}")
```

Inference

```
model.eval()  
context = torch.zeros((1, 1), dtype=torch.long) # Start token  
generated = model.generate(context, max_new_tokens=100, temperature=0.9)
```

Optimizing Training Speed

```
class OptimizedGPT2(GPT2):  
    """GPT-2 with optimizations for faster training"""
```

```
@torch.compile # PyTorch 2.0+ JIT compilation  
def forward_compiled(self, idx, targets=None):  
    return self.forward(idx, targets)
```

Enable mixed precision

```
from torch.cuda.amp import autocast, GradScaler

model = GPT2(config).to('cuda')
scaler = GradScaler()
optimizer = torch.optim.AdamW(model.parameters(), lr=6e-4)

for epoch in range(num_epochs):
    for X, Y in dataloader:
        X, Y = X.to('cuda'), Y.to('cuda')

        with autocast(dtype=torch.float16):
            logits, loss = model(X, Y)

            optimizer.zero_grad()
            scaler.scale(loss).backward()
            scaler.unscale_(optimizer)
            torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
            scaler.step(optimizer)
            scaler.update()
```

Use Flash Attention for faster attention computation

Requires: pip install flash-attn

```
from flash_attn import flash_attn_func
```

In CausalSelfAttention.forward:

```
attn_output = flash_attn_func(q, k, v,
causal=True)
```

Evaluation on GPT-2 Benchmarks

```
def evaluate_on_hellaswag(model, tokenizer, num_examples=100):
    """Evaluate on HellaSwag benchmark (reading comprehension)"""
```

```
from datasets import load_dataset

dataset = load_dataset("hellaswag")

correct = 0

for i, example in enumerate(dataset['test'][:num_examples]):
    context = example['ctx']
    endings = example['endings']
    label = int(example['label'])

    best_score = float('-inf')
    best_choice = 0

    for j, ending in enumerate(endings):
        full_text = context + " " + ending
        tokens = tokenizer.encode(full_text)
        tokens = torch.tensor(tokens).unsqueeze(0)

        with torch.no_grad():
            logits, _ = model(tokens)

        # Score: average log probability
        probs = F.softmax(logits[0, -len(ending.split()):], dim=-1)
        score = torch.log(probs + 1e-8).sum().item()

        if score > best_score:
            best_score = score
            best_choice = j

    if best_choice == label:
        correct += 1

accuracy = correct / num_examples
```

```
print(f"HellaSwag accuracy: {accuracy:.3f}")
return accuracy
```

Run evaluation

```
accuracy = evaluate_on_hellaswag(model, tokenizer)
```

11. Production LLM Systems

Scaling Considerations

Model parallelism for very large models

```
from torch.nn.parallel import DataParallel, DistributedDataParallel
```

Single node, multiple GPUs: DataParallel

```
model = DataParallel(model, device_ids=[0, 1, 2, 3])
```

Multi-node, distributed training: DistributedDataParallel

```
model = DistributedDataParallel(model, device_ids=[rank])
```

For models > 40B parameters, use tensor parallelism

(model layers split across GPUs)

```
from megatron.distributed import initialize_megatron
```

Or use inference engines that handle it:

- vLLM (recommended)

- Text Generation Inference (TGI)
- DeepSpeed Inference

Monitoring and Observability

```
import logging
from torch.utils.tensorboard import SummaryWriter
```

Setup logging

```
logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(name)
```

TensorBoard for metrics

```
writer = SummaryWriter()

def log_metrics(step, loss, lr, grad_norm):
    """Log training metrics"""
    logger.info(f"Step {step}: loss={loss:.4f}, lr={lr:.6f}, grad_norm={grad_norm:.2f}")
```

```
    writer.add_scalar('Loss/train', loss, step)
    writer.add_scalar('Learning_rate', lr, step)
    writer.add_scalar('Gradient_norm', grad_norm, step)
```

Usage in training loop

```
for step, (X, Y) in enumerate(dataloader):
    logits, loss = model(X, Y)
```

```
    optimizer.zero_grad()
    loss.backward()
    grad_norm = torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
    optimizer.step()

    if step % 100 == 0:
        lr = optimizer.param_groups[0]['lr']
        log_metrics(step, loss.item(), lr, grad_norm)
```

Conclusion

This comprehensive guide has covered:

1. **Foundations:** NLP, embeddings, neural networks, optimization
2. **Architecture:** Transformers, attention, position encoding
3. **Pre-training:** Autoregressive and masked language modeling
4. **Fine-tuning:** Full tuning, LoRA, QLoRA, instruction tuning, RLHF
5. **Prompting:** Zero/few-shot, CoT, ToT techniques
6. **RAG & Agents:** Knowledge retrieval, tool use
7. **Evaluation:** Metrics, benchmarking, safety
8. **Deployment:** Inference optimization, serving
9. **Custom Training:** Complete pipeline from data to inference
10. **GPT-2 Replication:** Understanding architecture deeply
11. **Production:** Scaling, monitoring, optimization

Next steps:

- Start with small experiments (GPT-2 small, 124M params)
- Scale up gradually as you understand the concepts
- Use modern frameworks: HuggingFace, LangChain, vLLM
- Contribute to open-source LLM projects
- Stay updated with latest techniques (Vision Transformers, Multimodal models, etc.)

The field of LLMs evolves rapidly. The fundamentals outlined here will remain relevant, but always check latest research papers and community practices.